

PracticalNotebook3

February 15, 2023

1 Practical Lab 3

In this lab, we will focus more on using machine learning models to solve mock real-world problems. Scikit-learn is often the go-to library for using traditional statistical machine learning models. Scikit learn provides a collection of data pre-processing and machine learning models that can seamlessly be integrated with other Python packages.

In the previous practical lab, we described the libraries Numpy, Pandas, and Matplotlib, which are useful for loading, transforming, and plotting data.

At the end of this lab, you should be more familiar with searching for documentation for library functions online, using linear and logistic regression function with Scikit-learn to solve a given problem and identify when to use which model. You should also be even more familiar with the libraries Numpy, Matplotlib, and Pandas and how you use these libraries in conjunction.

This lab will deal with:

- How to use and train a linear regression model
- How to use and train a logistic regression model.
- Some approaches on how you can investigate a new dataset, its most important features, distribution, etc.
- How to create a predictive model to predict the class or label of new unseen data.
- Metrics and visualization methods for evaluating a trained predictive model.

1.1 Machine learning

Scikit-learn is a (near) all-encompassing python library for machine learning.

When we talk about machine learning, it is crucial to distinguish between the types of problems we are trying to solve. By grouping and categorizing the kinds of problems we are trying to solve, we can identify commonalities between them and develop methods for solving similar tasks.

Having good data is essential for all machine learning models. If the data is incorrect, unfiltered, missing values, and so forth, then the model built from that data will perform poorly! The adage goes, “Garbage in, garbage out.”

Below is a simplified image of how different tasks are grouped for machine learning

1.1.1 Regression vs Classification

In this lab, we will mainly deal with supervised learning tasks. This means that we have data with a corresponding label. A supervised model aims to extract features for the data and learn to

predict the labels on unseen data as accurately as possible.

In regression and classification tasks, the goal is to find a function, $f(X)$, which is the most likely to have generated the data we observe. The major difference between a regression and classification task is:

1. For **regression**, the label for the data is some number and is continuous. It may range from anywhere between $-\infty$ to ∞ . It can take any value on the real number line.

Example tasks:

- Predicting house prices based on location and other factors.
- Predicting the bpm for a song based on its features.

2. For **classification**, the label is some category, class, or other groupings. Based on which features the data has, one wishes to group/classify the data's correct type. **Example tasks:**

- Predicting cancer cells for images of skin lesions
- Predicting if the video feed for a self-driving car is your side of the road or not.
- Predicting the most likely score of movies, based on what you previously have liked.

2 a) Regression Task

2.1 House Prices

This section will describe a linear regression task by looking at a common dataset, called the Boston housing dataset. Our goal is to predict what the price of any house would be given certain attributes but to do that, it is a good starting point to first inspect the data!

```
[ ]: from sklearn import datasets
      from sklearn.datasets import fetch_california_housing

      california = fetch_california_housing()
      print("Columns of the dataset:\n", california.keys())

      print("\nShape of the data", california.data.shape)
      print("Number of examples", california.target.shape)
```

Columns of the dataset:

```
dict_keys(['data', 'target', 'frame', 'target_names', 'feature_names',
'DESCR'])
```

Shape of the data (20640, 8)

Number of examples (20640,)

```
[ ]: print("\nShape of the data", california.feature_names)
```

Shape of the data ['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude']

For the Scikit-learn datasets, we can read more about a particular dataset by printing its description:

```
[ ]: print(california.DESCR)
```

```
.. _california_housing_dataset:
```

```
California Housing dataset
```

```
-----
```

```
**Data Set Characteristics:**
```

```
:Number of Instances: 20640
```

```
:Number of Attributes: 8 numeric, predictive attributes and the target
```

```
:Attribute Information:
```

```
  - MedInc           median income in block
  - HouseAge         median house age in block
  - AveRooms         average number of rooms
  - AveBedrms        average number of bedrooms
  - Population       block population
  - AveOccup         average house occupancy
  - Latitude         house block latitude
  - Longitude        house block longitude
```

```
:Missing Attribute Values: None
```

This dataset was obtained from the StatLib repository.

<http://lib.stat.cmu.edu/datasets/>

The target variable is the median house value for California districts.

This dataset was derived from the 1990 U.S. census, using one row per census block group. A block group is the smallest geographical unit for which the U.S. Census Bureau publishes sample data (a block group typically has a population of 600 to 3,000 people).

It can be downloaded/loaded using the

```
:func:`sklearn.datasets.fetch_california_housing` function.
```

```
.. topic:: References
```


```
  - Pace, R. Kelley and Ronald Barry, Sparse Spatial Autoregressions,
    Statistics and Probability Letters, 33 (1997) 291-297
```

We can visually inspect the dataset to get a sense of it.

```
[ ]: import matplotlib.pyplot as plt

plt.hist(california.target, bins=10, rwidth=0.9, alpha=0.8)
plt.xlabel("Price ($100K)")
plt.ylabel("Count")
```

```
[ ]: Text(0, 0.5, 'Count')
```



Once we have gained better oversight of the data, we will want to fit a linear regression model on the data. Once we have generated our predictive model, we can use it to make predictions of what new housing prices are likely to be.

ASSIGNMENT a) a) Look up the documentation for importing a linear regression model from Scikit-learn.

b) Use it to train a linear regression model on the Boston house prices dataset.


- c) Use the trained regression model on the dataset to predict what the housing prices should be from the original dataset. Is it accurate?
- d) Plot the prediction your model made against the true labels with a scatter plot.

```
[ ]: import matplotlib.pyplot as plt
price = fetch_california_housing()
# YOUR CODE HERE
#
#
from sklearn.linear_model import LinearRegression

x = price.data
y = price.target

reg = LinearRegression().fit(x, y)
prediction = reg.predict(x)

# Code provided by us to compare the generated predictions against the actual
↪ values
xlim, ylim = [0, 5], [0, 5]
plt.scatter(y, prediction)
plt.plot(xlim, ylim, "--k")
plt.xlabel("True price ($1K)")
plt.ylabel("Predicted price ($1K)")
plt.show()
```



2.1.1 Train-Test split

In the above example, you saw how you could create a predictive model to fit all of the data. However, this approach has a problem: If your model overfits the data, you have no way of knowing since the model has been trained to work the best for that dataset.

Instead, we split the data into a training and testing dataset or even a train, test, and validation dataset. This way, we can use the training dataset for training the model and use the leftover unseen data (test set) to evaluate how well our model performs on new data. A naive way would be to split it based on indexing, say the first 80 samples are training data, and the test is testing. But, there is a high likelihood that samples next to each other were collected with high correlation. Maybe a dataset is measurements of temperatures taken every week. If we split based on indexing, our model might only train on temperatures for spring and summer!

Therefore, a better method is to choose which data points are in the training and testing set randomly. This approach is sometimes called a train-test split - and when splitting the data, the training dataset should contain more samples than the testing set since ML models require a lot of data to be trained well.

ASSIGNMENT b) a) Train a Linear Regression model again, but this time split the dataset into a training and testing dataset. Either of these splits is allowed. The first number indicates the percentage of the whole data that is in the training set, and the other, the percentage of data in the testing set:

- (70/30)
- (75/25)
- (80/20)
- (85/15)
- (90,10)

b) Use a Linear Regression model and make predictions on the **test dataset**, data which the model has not trained on. Save the predicted labels generated for the model in a variable named y_{pred} .

c) Create the same scatter plot between the predicted target, y_{pred} and the actual target, y_{test} as you did before in the previous section of the assignment. This plot will illustrate how close your predictive model is with its predictions compared to the actual labels.

d) Create a title for your plot stating how you chose to split the training and testing data. The plot's title should also display the *coefficient of determination* R^2 - i.e., how well the model's predicted labels compares to the true labels.

Example of how the title of the plot may look like:

“Housing Prices, split (80/20) - R2: 0.71”

(optional): Test different training and testing splits. How does the *coefficient of determination* for your predictive model change?

```
[ ]: import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.datasets import fetch_california_housing
california = fetch_california_housing()

# YOUR CODE HERE
#
#
from sklearn.model_selection import train_test_split

split = 0.2

x = california["data"]
y = california.target

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=split)

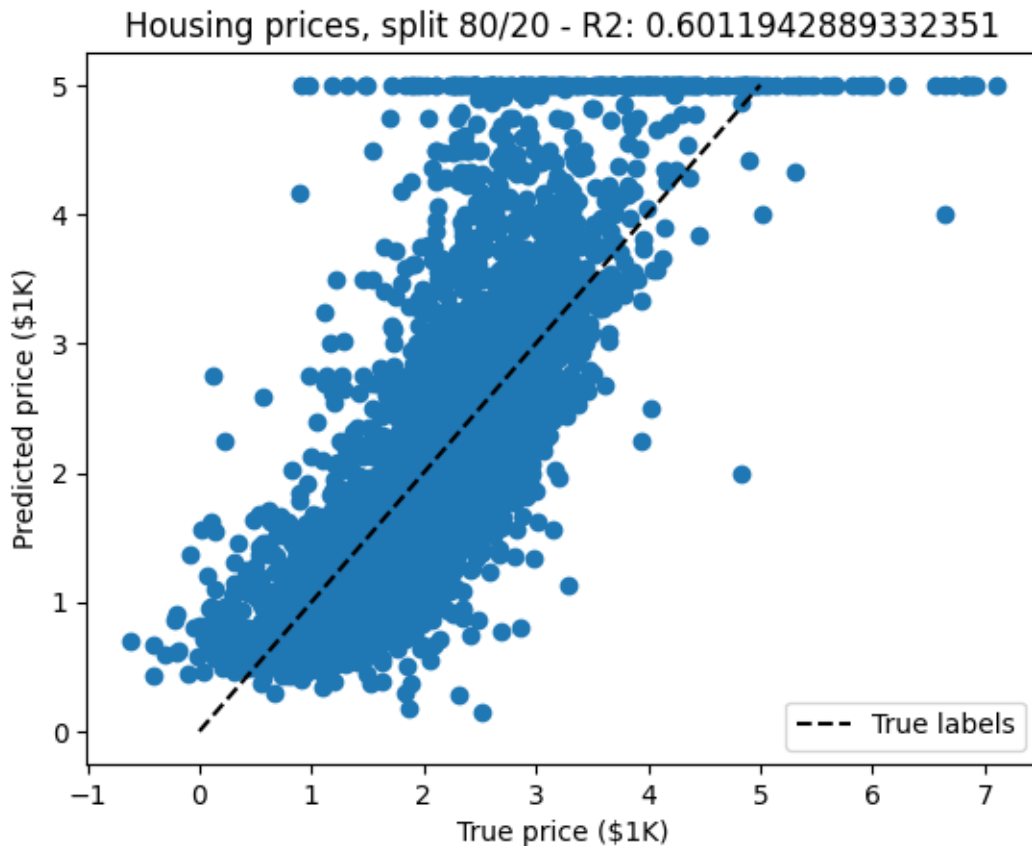
reg = LinearRegression().fit(x_train, y_train)
prediction = reg.predict(x_test)

# CODE PROVIDED TO COMPARE YOUR RESULTS AGAINST THE ACTUAL LABELS
```

```

xlim = [0, 5]
ylim = [0, 5]
plt.plot(xlim, ylim, '--k', label=f"True labels")
plt.scatter(prediction, y_test)
plt.title("Housing prices, split " + str(int(100*(1-split))) + "/" +
    ↪str(int(100*split)) + " - R2: " + str(reg.score(x_test, y_test)))
plt.xlabel("True price ($1K)")
plt.ylabel("Predicted price ($1K)")
plt.legend()
plt.show()

```



This type of visual measurement is a quick and effective way to verifying that the trained model behaves as expected. However, to compare several different models, we need a more robust and mathematical precise measure. We will discuss some of those methods in the next section.

3 b) Classification task

3.0.1 The iris dataset

A commonly used dataset for testing classification algorithms is the so-called “Iris dataset.” It consists of measurements from three different types of Irises:

- a) Iris-Setosa
- b) iris-Versicolor
- c) Iris-Virginica

From these are measured features of the different flowers, such as:

- sepal length (cm)
- sepal width (cm)
- petal length (cm) -petal width (cm)

Based on the features of these different flowers (measurements), we want to predict which is the most likely Isis (types of flowers) given these measurements?

```
[ ]: import matplotlib.pyplot as plt
import numpy as np

from sklearn.datasets import load_iris
iris = load_iris()

# We can view the columns of this dataset
print("Iris dataset cloumnts: \n", iris.keys())

# View the features for the dataset
iris_features = iris['feature_names']
print("\nThe Iris dataset features, ", len(iris_features), ":\n", iris_features)
```

Iris dataset cloumnts:

```
dict_keys(['data', 'target', 'frame', 'target_names', 'DESCR', 'feature_names',
'filename'])
```

The Iris dataset features, 4 :

```
['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width
(cm)']
```

Each sample will also have a target. In this case, the class (types of Irises) the sample belongs to.

```
[ ]: iris_targets = iris.target
iris_classes = list(set(iris_targets)) # Get only the unique targets
iris_target_names = ["Setosa", "Versicolor", "Virginica"]

print("Iris targets and names:\n" + 20*"_" + "\n")
for i, name in zip(iris_classes, iris_target_names):
    print(f"{i}: Iris-{name}")
```

Iris targets and names:

```
-----  
0: Iris-Setosa  
1: Iris-Versicolor  
2: Iris-Virginica
```

```
[ ]: print("Example from the Iris dataset:\n Sample =",iris['data'][0])  
      print("\nThis dataset has {} features".format(len(iris['data'][0])))
```

```
Example from the Iris dataset:  
Sample = [5.1 3.5 1.4 0.2]
```

```
This dataset has 4 features
```

```
[ ]: nr_samples, nr_features = iris.data.shape  
      print("Number of samples:", nr_samples)  
      print("Number of features:", nr_features)
```

```
Number of samples: 150  
Number of features: 4
```

As we can see, each measurement (feature) of the data adds additional dimensions. To investigate the types of differences between the datasets, we may want to examine if we can see clear distinctions between them based on their features. However, since there are four features, we either need to choose between two attributes and compare them or reduce the dimensions (features) down to something we can more easily visualize. The latter is a collection of dimensionality reduction methods and will be described in an upcoming lab.

We can visualize one or two features simultaneously via a scatterplot, histogram, or similar methods.

Assignment c)

- Plot a histogram for each of the features of the Iris dataset - ie, four histograms
- Determine the feature that could best categorize the type of Iris. How can you find which part is the most suitable for it?

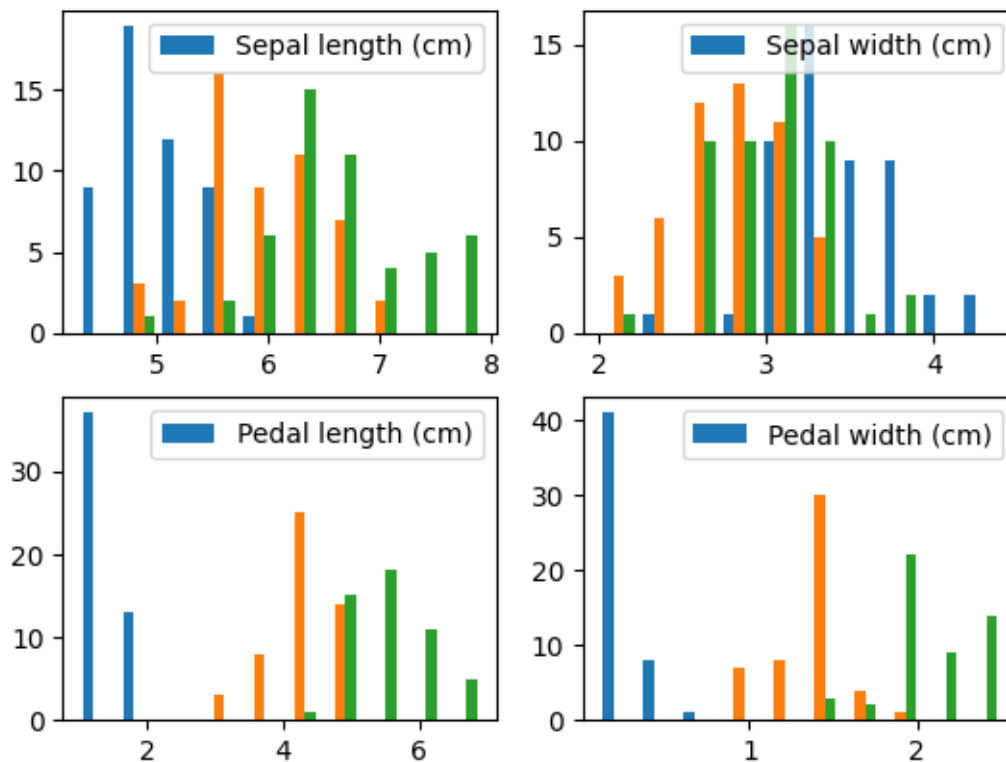
```
[ ]: import matplotlib.pyplot as plt  
      iris_target_names = ["Setosa", "Versicolor", "Virginica"]  
  
      # YOUR CODE HERE  
      #  
      #  
  
      sep_len = iris.data[0:50, 0], iris.data[50:100, 0], iris.data[100:150, 0]  
      sep_wid = iris.data[0:50, 1], iris.data[50:100, 1], iris.data[100:150, 1]  
      ped_len = iris.data[0:50, 2], iris.data[50:100, 2], iris.data[100:150, 2]  
      ped_wid = iris.data[0:50, 3], iris.data[50:100, 3], iris.data[100:150, 3]
```

```
fig, ax = plt.subplots(2, 2)
fig.suptitle("Iris Data Histograms")

ax[0, 0].hist(sep_len, label="Sepal length (cm)")
ax[0, 0].legend()
ax[0, 1].hist(sep_wid, label="Sepal width (cm)")
ax[0, 1].legend()
ax[1, 0].hist(ped_len, label="Pedal length (cm)")
ax[1, 0].legend()
ax[1, 1].hist(ped_wid, label="Pedal width (cm)")
ax[1, 1].legend()
plt.show()
```

"""
 YOUR ANSWER FOR THE BEST FEATURE (IN YOUR OPINION) BY NAME HERE:
 Pedal length and pedal width are the best features in my opinion. They seem to
 be the most distinct between the three classes.
 """

Iris Data Histograms



```
[ ]: '\nYOUR ANSWER FOR THE BEST FEATURE (IN YOUR OPINION) BY NAME HERE:\nPedal length and pedal width are the best features in my opinion. They seem to be the most distinct between the three classes.\n'
```

From the histogram, you likely found one or possibly more features that you thought can best describe/separate the data. To achieve an even better separation between the labels, we can use additional features.

Assignment d) We now want to create a scatter plot using the two most important features to best separate the Iris dataset into its distinct groups based on the type of Iris. Your task is to:


- Create a scatter plot between different features and find the two you think are the best at separating data into distinct groups.
- Present your result below by labeling the Iris type by name and which color is assigned in the plot.
- The plot should look something like this, but you may find better splits:
- (Optional) if you feel comfortable with Python, you are allowed to use other plotting libraries to solve this task if you find it more convenient.

```
[ ]: # YOUR CODE HERE
#
# Pedal length and sepal width
setosa = iris.data[0:50, 2:4]
setosa_length = setosa[:, 0]
setosa_width = setosa[:, 1]

versicolor = iris.data[50:100, 2:4]
versicolor_length = versicolor[:, 0]
versicolor_width = versicolor[:, 1]

virginica = iris.data[100:150, 2:4]
virginica_length = virginica[:, 0]
virginica_width = virginica[:, 1]

plt.scatter(setosa_length, setosa_width, label="Setosa")
plt.scatter(versicolor_length, versicolor_width, label="Versicolor")
plt.scatter(virginica_length, virginica_width, label="Virginica")
plt.legend()
plt.xlabel("Pedal length (cm)")
plt.ylabel("Pedal width (cm)")
plt.show()
```



ASSIGNMENT e)

- Plot the correlation plot between all the different features.
- In WRITING also explain BRIEFLY what a correlation plot means
- Do the correlation plot's results match what you found to be the best features for separating the data into labels?
- **Hint:** You may want to convert the data into a DataFrame to more easily make a correlation plot

```
[ ]: # YOUR CODE HERE
#
#
import pandas as pd
import seaborn as sn
iris_data = iris.data

df = pd.DataFrame(iris_data, columns=iris_features)
correlation = df.corr()
sn.heatmap(correlation, annot=True)
plt.show()
"""
```

YOUR BRIEF ANSWER TO WHAT A CORRELATION PLOT MEANS HERE:

The plot shows the correlation between the different features. So if the value is close to 1, the features are highly correlated. If the value is close to 0, the features are not correlated. If the value is close to -1, the features are negatively correlated.

For the best separation between the classes, I would choose pedal length and pedal width. They have the highest correlation, which leads to better separation

"""



```
[ ]: '\nYOUR BRIEF ANSWER TO WHAT A CORRELATION PLOT MEANS HERE:\nThe plot shows the correlation between the different features. So if the value is close to 1, the features are highly correlated. If the value is close to 0, the features are not correlated. If the value is close to -1, the features are negatively correlated.\nFor the best separation between the classes, I would choose pedal length and pedal width. They have the highest correlation, which leads to better separation\n'
```

ASSIGNMENT f)

- Create a logistic regression classifier to predict new datapoint into correct Iris labels
- Use a train test split of (80/20)
- Compute the accuracy of the fitted machine learning model and call it acc

```
[ ]: # YOUR CODE HERE
#
#
from sklearn.linear_model import LogisticRegression
x = iris.data[:, 2:4]
y = iris.target
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2)

model = LogisticRegression().fit(x_train, y_train)

acc = model.score(x_test, y_test)

# CODE PROVIDED TO TEST YOUR FITTED MODEL
assert acc >= 0.95, "You haven't reached a high enough accuracy."
```


The final part, which only works for classification tasks, is for us to investigate how well the classifier faired when predicting the new labels. A handy method is a confusion matrix. It compares how many of the actual labels were classified correctly and (importantly) which labels the model made incorrect predictions.

ASSIGNMENT g)

- Create a confusion matrix to illustrate how the labels of the actual Iris dataset differ from the predicted labels

```
[ ]: # YOUR CODE HERE
#
#
from sklearn import metrics

prediction = model.predict(x_test)
actual = y_test
confusion = metrics.confusion_matrix(actual, prediction)
cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix=confusion,
↪display_labels = iris.target_names)
cm_display.plot()
plt.show()
```



4 Assignment h)

In an old version of this netbook we used the Boston housing dataset. If you try to import it then you get an error message.

```
[ ]: from sklearn import datasets

price = datasets.load_boston()
```

Read the error message and look at article by Carlise. What do you think about the ethical implications of using such a dataset? In the cell below write at least 10 so lines about what you think of using such a dataset and the article by Carlise. You should change the cell type below to Markdown to write your relections.

5 Your reflections

The boston dataset seems to be veiled in controversy, due to many ethical issues. This seems to range from things like capping the housing costs at 50k, to recording the proportion of the population that is black in a certain area. This raises a lot of questions, about the validity of the dataset, as well as the ethical code that the original creators lived by. In this case, they even use the proportion of the population to suggest that areas with more people of color has lower housing

costs. This makes me feel as if this is not the sort of dataset that one would want to use. From a pure ethical perspective, it is quite daunting to see such a backwards idea make its way into a well used dataset. But even from a pure computer science perspective, the fact that this was added makes the data simply put wrong. So, the idea that adding this would be because of some logical reason seems to go out the window. All in all, it seems reasonable to stop using this dataset, as today there are far better, far larger and far more ethically sound datasets to use.

6 Where do we go from here?

6.0.1 Methods for improving model performance even more

We have now trained linear and logistic regression models on two different tasks. These models have had a reasonably good performance. However, we can improve it in several different ways, which we delve into in the upcoming practical notebooks. Some examples include:

1. Normalize and clean the data
2. Compare and use other models such as SVMs, Trees, KNNs, generalized linear models, neural networks, etc.
3. We can artificially generate more data samples. More training data yields a better performing model
4. We can allow the model to see more validation samples (cross-validation) artificially
5. Measure the model's ability when making predictions. When is it incorrect and in what way (Accuracy, Recall, Precision, F1, and ROC-curves)
6. Certain aspects are more important than others when evaluating a model. Therefore it might be useful to use different evaluation metrics for various tasks

[]: